

Week 2 Task Submission: Database & API Development

Name: Rishav Dev Tiwari

Program: Fusemachines AI Fellowship 2026

1. Project Overview

This project implements a fully functional REST API for the ClassicModels database utilizing a 4-layer clean architecture (Connection, Data Format, Logic, and API Layer). The application adheres to the Twelve-Factor App methodology, specifically focusing on Dependency Management, Config Management, Backing Services, and Concurrency.

2. Environment & Database Setup (Factor II, III, & IV)

The application utilizes Docker to ensure a consistent environment and to manage PostgreSQL as an isolated backing service. Database credentials and variables are isolated in a .env file.

The database is containerized using docker-compose.yml. Upon the initial startup, the container automatically executes seed.sql via the /docker-entrypoint-initdb.d/ volume bind, successfully building all 8 tables and seeding them with data.

```
PS D:\Fusemachines_Fellowship\Week2_SoftEng\project> docker-compose up -d
time="2026-05-12T23:44:17+05:45" level=warning msg="D:\Fusemachines_Fellowship\Week2_SoftEng\project\docker-compose.yml: the attribute `version` is obsolete, it will be ignored, please remove it to avoid potential confusion"
[*] up 18/18
✓ Image postgres:latest Pulled 89.1s
✓ Network project_default Created 0.1s
✓ Container postgres_db Started 4.2s
```

3. API Testing & Documentation

The API leverages FastAPI automatic interactive documentation (Swagger UI). Each of the 8 tables features its own tagged section containing GET and COUNT endpoints.

ClassicModels API

0.1.0

OAS 3.1

/openapi.json

Week 2 Assignment API

ProductLines



GET	/productlines/count	Get Productlines Count	▼
GET	/productlines/	Get Productlines	▼

Customers



GET	/customers/count	Get Customers Count	▼
GET	/customers/	Get Customers	▼

Dashboard



GET	/overall_counts	Get Overall Counts	▼
-----	-----------------	--------------------	---

Dashboard



GET

/overall_counts

Get Overall Counts

^

Parameters

Cancel

No parameters

Execute

Clear

Responses

Curl

```
curl -X 'GET' \
  'http://localhost:8000/overall_counts' \
  -H 'accept: application/json'
```

Request URL

```
http://localhost:8000/overall_counts
```

Server response

Response headers

```

content-length: 21
content-type: text/plain; charset=utf-8
date: Tue, 12 May 2026 18:02:41 GMT
server: uvicorn

```

Responses

Code	Description	Links
200	Successful Response	No links

Media type

application/json ▼

Controls Accept header.

Example Value | Schema

```

"string"

```

4. Concurrency Implementation (Factor VIII)

To achieve high performance under load, the application utilizes asynchronous concurrent processing for the master dashboard endpoint (/overall_counts).

Instead of querying the 8 database tables sequentially (which creates blocking delays), the API uses `asyncio.gather()` combined with `asyncio.to_thread()`. This allows the FastAPI server to dispatch all 8 `SELECT COUNT(*)` queries to the PostgreSQL database simultaneously, drastically reducing the total response time.

```

INFO:     Uvicorn running on http://127.0.0.1:8000 (Press CTRL+C to quit)
INFO:     Started reloader process [10368] using StatReload
INFO:     Started server process [34300]
INFO:     Waiting for application startup.
INFO:     Application startup complete.
INFO:     127.0.0.1:60425 - "GET /docs HTTP/1.1" 200 OK
INFO:     127.0.0.1:60425 - "GET /openapi.json HTTP/1.1" 200 OK
2026-05-12 23:47:42,036 - api_logger - INFO - Incoming request: GET /overall_counts
2026-05-12 23:47:42,036 - api_logger - INFO - Starting concurrent tasks...
2026-05-12 23:47:42,045 - api_logger - INFO - Executing query: SELECT COUNT(*) FROM customers

```

Github Repository: <https://github.com/rishavdevtiwari/fusemachines-fellowship.git>